

- Today, most of the compute & storage intensive work goes on in networked servers.
- There is hardly any work on a large scale that happens on local computers.
- Everything is moving to the cloud.
- All large companies have their own datacenters for processing data, or use public cloud offerings.
- All the data present in public clouds like Amazon AWS, Google Compute Engine or Microsoft Azure, resides on VMs. Virtualization is in the heart of public cloud offerings.
- In this class, we will look at the virtualization offered by the public cloud providers, such as Google & Amazon.

- We will see what these companies offer in terms of VMs, containers, serverless, etc.
  - Then we will look at how major companies like netflix / twitter / uber use the services for all of their data processing.
  - This class will serve as a motivation for the other classes, & will talk about the usefulness & impact of virtualization, before looking into "how" to virtualize in the remainder of the semester.
- 

### Timeline :

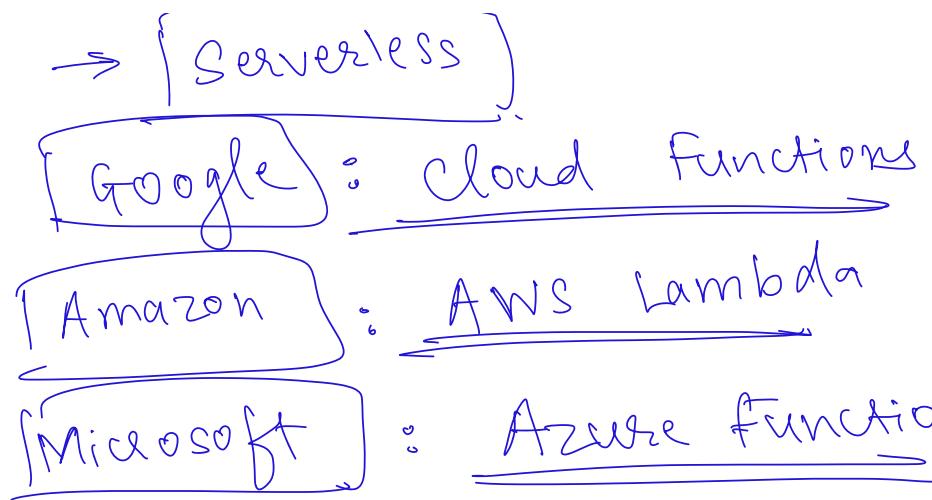
- ① VMs.
- ② Containers.
- ③ Unikernels
- ④ Serverless
- ⑤ New VM technologies.

Twitter ad-analytic framework:  
managed by Google cloud.

- GCP:
  - Spot VMs (GCP)
  - can be used to create VMs that are fault tolerant.
  - The VMs are created on-demand. They are allocated resources when there are available idle resources.
  - If the load on servers becomes high, the spot VMs are either paused or killed.
  - User applications running on spot VMs should be fault tolerant
  - Not good if there are tight SLAs. The fail latency becomes high due to spot VMs.
  - 91% lower cost than general-purpose VMs.
  - Used by maps, YouTube, etc

- Harvest VMs. (Azure)
  - More flexible than Spot VMs.
  - The resources in Harvest VMs grow & shrink according to the unallocated resources in underlying physical machine
  - Does not get killed. Only killed if the resources go below the minimum required resources of the VM.
  - More utilization of data centers
- Burstable VMs.
  - VMs that are allocated some minimum amount of resources, but that can get more resources - on-demand in a best-effort manner if the application load increases.





Functions that run in the cloud offer FaaS. You just need to write a function, such as a serverless web app. Then it is associated with an API. The function gets called on the API, & runs in the cloud.

Example:

Coca-Cola used AWS Lambda

use QR-code on phone to scan the coca booth. Automatically shows UI with options. On click, the AWS Lambda function is executed, & drink is poured in the cup.

Only pay for the invocation of function.  
Payment happens at millisecond granularity.

## Disadvantage :

No state. It needs to be stored in separate db outside of application.

Function needs to be self-sufficient in the data it uses. Cannot rely on global variables, or data of other functions. Hard in some cases.

## Durable Functions : Stateful Serverless

### functions

In some scenarios, Azure manages to hold state for your functions & hence makes it easy for us to use serverless in real world.

VMs

GCE: 2012

Amazon EC2 : 2006

Microsoft Azure : 2008

Charged on per-minute usage.

Containers

Kubernetes: 2014

Docker: 2013

GKE charged in \$0.1 per  
cluster per hour.

Serverless

2017

: Google  
cloud functions.

- Developed in company.
  - ↳ - Job: long period of time.
- pre-virtualization.  
(so yrs ago.)
  - competing something
  - How?
  - either buy lot of machine.
  - Amazon EC2, Azure, etc  
↳ with price
  - only billed while time is running.
- Job: not all the time. Only get erratically, 10 times a month.
  - EC-2. Big cost of keeping it running
  - Serverless. Info. from external sources.  
Bill lower.
- Containers: Developer working in company. Just send depending.

1st jump: machine. single machine  
→ on-prem machine. V.M.  
→ on-prem. V.M.  
→ cloud  
Each thing. : How people did  
before  
then.  
Connect to parts of class.

- 
- Last time: history of Virtualization.
  - We talked about timeline of when VMs introduced, containers introduced & serverless introduced
  - Today: motivate virtualization, look at real-world examples & impact of virtualization.
  - Imagine that we have an IT company
    - Accounting dept
    - Developers - backend
    - Marketing dept.
    - Developers - front-end.

- In 90s, there were no machines.
- Every employee: one computer with own monitor, keyboard.
- ↳ Each computer: \$1000 - \$2000 minimum.
- IT support staff: manage SW on computers, update O.S., etc.
- Initial startup investment: very high.  
↳ 10s of 1000s of \$.

### I<sup>st</sup> Jump of Virtualization.

- on-prem VMs.
- Instead of separate desktop, server room with all machines.
- VMs for every employee depending on requirements.
- Underlying physical machine: manage CPU, mem., storage & networking across VMs.
- ↳ Every employee didn't need a separate server.
- ↳ However, IT support staff was still needed, to support the

- Half the servers needed
- ↳ inventory cost.

Look at when we learn about VMs next week

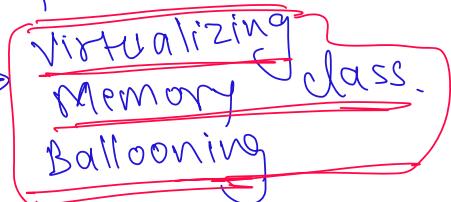
### I<sup>nd</sup> Jump → CLOUD.

- There was no physical resource needed.
  - = simply launch a VM in the cloud, using Amazon EC2, Google Cloud Platform or Microsoft Azure.
  - No IT support staff needed.
  - Very low initial investment on resources, low overhead to start a company almost everything must run in a virtualized setting.
- Financial year-end = Calculate profits & revenue.
- Run the software for multiple days & do lot of computation.
  - Because of 2nd jump of virtualization:
    - 16 vCPUs, 32 GB mem, 200 GB SSD: \$250/month.
    - or - \$8.3/day.
    - Launch VMs, costing total be done with it, instead of paying & maintaining for own machines.
    - +3 medium: 2 vCPUs, 6 GB mem, 100 GB SSD: \$2/month
    - or - \$6/month.
  - You get paid for every second use of the VM, & it is maintained & managed by

## Cloud providers.

- In fact to reduce costs even further I tailor to our needs different VMs are present.

- ① spot VMs
- ② Hail VMs
- ③ Burstable VMs.



IT team comes up with a complex software product, that depends on a number of packages & their versions, such as python 3, JDK v7.2 → 100MB of space  
→ Old days: no virtualization: ship 500 MB DRAM  
the package itself sell the software DRAM

- customer machine different, & would fail.
- days of debugging required.
- The software is computation heavy, but requires only about 100 MB of memory.
- ship it in VM, of 10s of GBs.
- slow to start-up, too heavy.
- Any bug, re-install VM at the client end.
- CONTAINERS

- package the software & dependencies, ship to client
- client can run at their end.
- Any crash, same environment & easily reproducible
- Container Orchestration

framework:

Google Kubernetes Engine.

Platform-as-a-service.

- Put container & orchestrate with GKE. It places your containers & executes in the cloud on its own
- you don't need to worry about VMs or anything.
- You only mention the # CPUs, mem. & storage needed for app. GKE manages the rest.
- We will learn about containers after JMS.

## Netflix

- - - - -
- Come up with a software that is only called erratically. probably 10 times a month, it gets invoked.

e.g. V

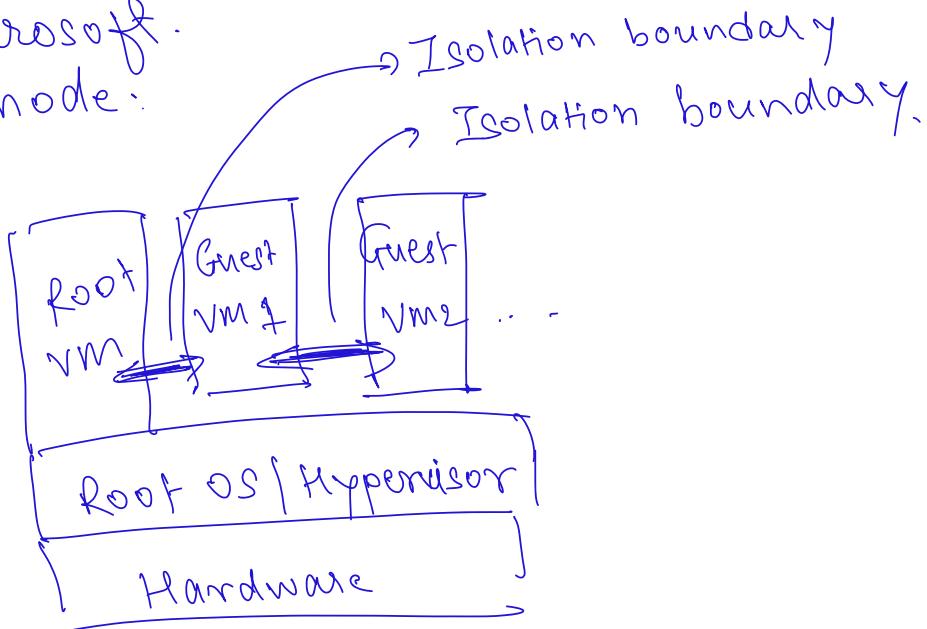
- Having a VM running all the time causes you to pay for the VM all year.
  - Same with container.
- 
- Serverless: write the service, stores data in Amazon S3.  
What is Amazon S3.
    - Write an API for the service.
    - Push it to AWS Lambda.
  - Whenever API is called, it is run on AWS Lambda.
  - No VM
  - No container.
  - Function-as-a-service.
  - \$1.37/month only
- ↳ Coca-Cola
- We will learn about Serverless towards the end

↳ Coca-Cola free-style.

## MICROSOFT AZURE

Uses hypervisor named Hyper-V by Microsoft.

- Each node:



### Cross VM attack:

- When victim VM & adversary-controlled VM are on the same host, & adversary-controlled VM snoops or causes slowdown of victim VM.
  - Side-channel attacks

Fabric Controller → Sophisticated placement of VMs on hosts. VMs don't know other co-existing VMs on the same host.

- Root VM
  - ↳ Special VM that runs special OS

## AWS Nitro

September 2020.

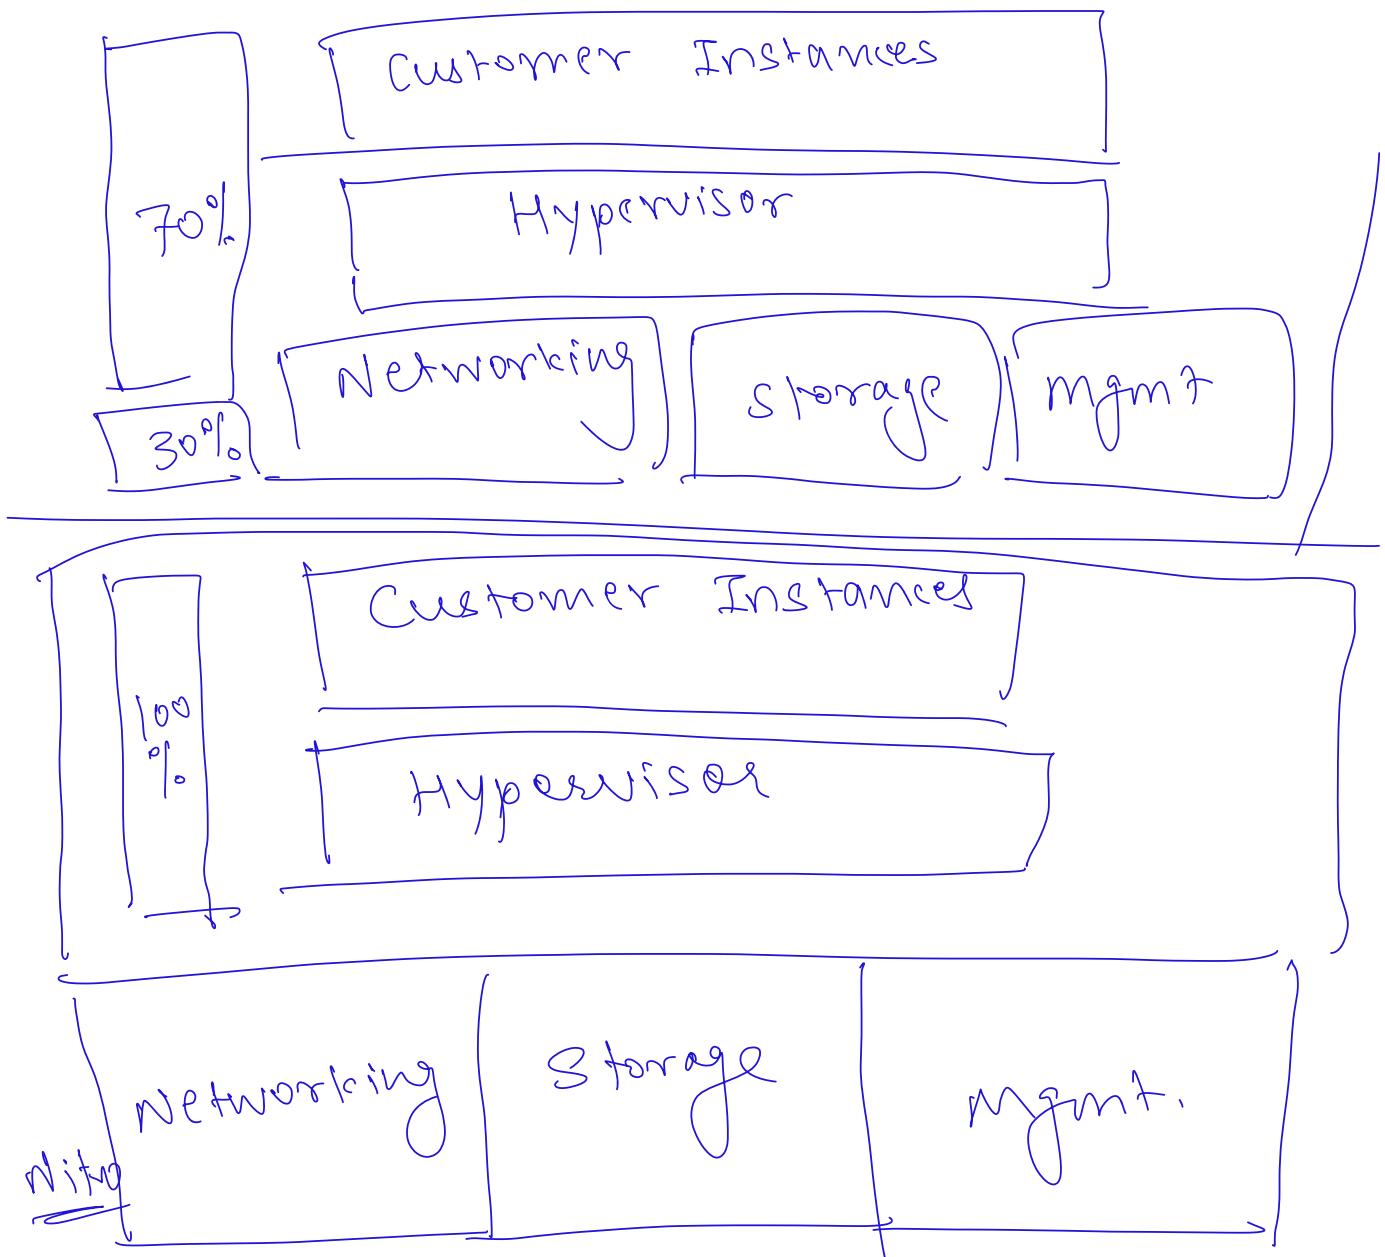
offload virtualization functions to dedicated hardware & software.

In early days of EC2, Xen hypervisor was used, & was modified to support multi-tenancy, to increase utilization & reduce costs.

However, with Xen, as much as 30% of the resources in an instance were allocated to hypervisor.

Nitro components :-

- ① Nitro cards : Accelerated IO for functions, such as instance storage.
- ② Nitro Hypervisor : Complete hypervisor offloaded to hardware.  
Nitro enabled Amazon to make hypervisor layer optional & offer baremetal instances.



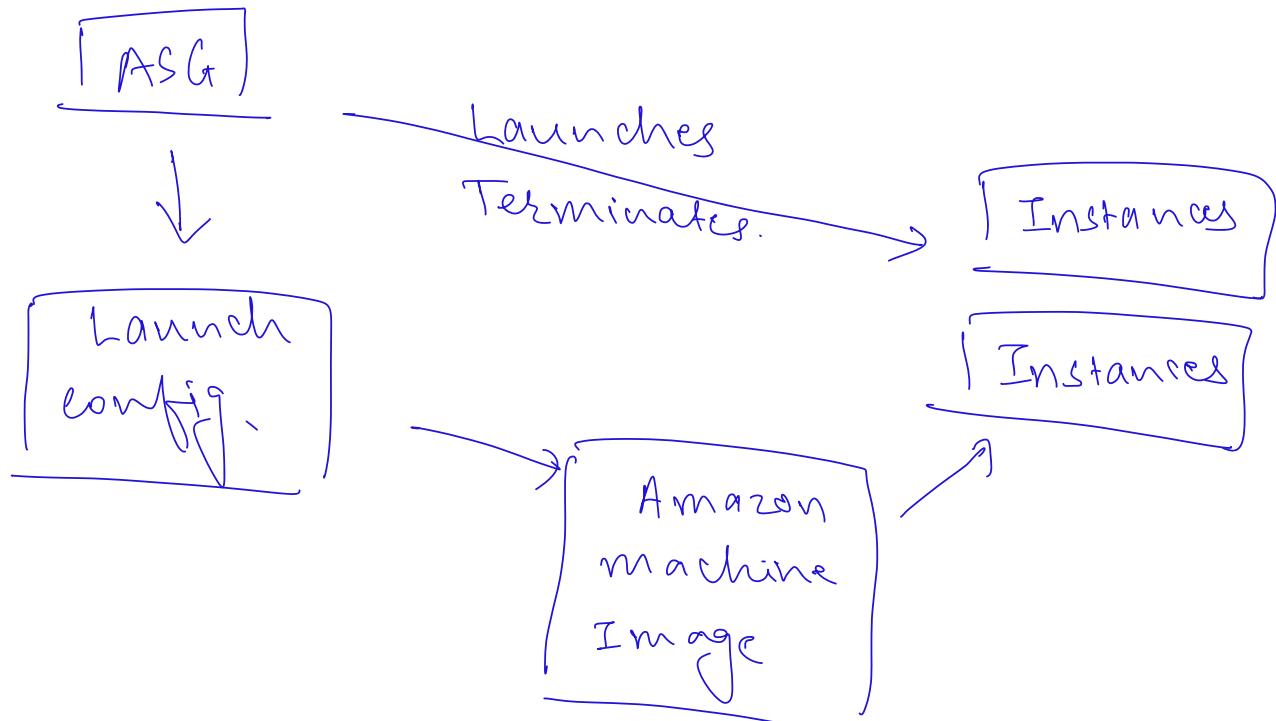
Faster Networking : 100 Gb/s.

Faster storage : 4x faster.

Security : Hardware enforced.

Cheaper : Better utilization of resources.

Before



After

